

CREATED BY

UDAYABHANU NAYAK



```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: df = pd.read_csv(r'C:\Users\Udayabhanu\Downloads\Stroke Prediction\Stroke Data
```

```
In [3]: df.head()
```

Out[3]:

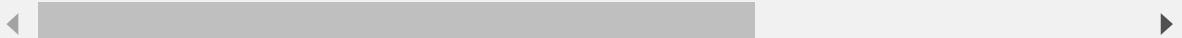
	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type
0	9046	Male	67.0	0	1	Yes	Private	Urban
1	51676	Female	61.0	0	0	Yes	Self-employed	Rural
2	31112	Male	80.0	0	1	Yes	Private	Rural
3	60182	Female	49.0	0	0	Yes	Private	Urban
4	1665	Female	79.0	1	0	Yes	Self-employed	Rural



```
In [4]: df.tail()
```

Out[4]:

	id	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_ty
5105	18234	Female	80.0	1	0	Yes	Private	Urb.
5106	44873	Female	81.0	0	0	Yes	Self-employed	Urb.
5107	19723	Female	35.0	0	0	Yes	Self-employed	Ru
5108	37544	Male	51.0	0	0	Yes	Private	Ru
5109	44679	Female	44.0	0	0	Yes	Govt_job	Urb.



In [5]:

```
df = df.drop('id', axis=1)
df
```

Out[5]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg
0	Male	67.0	0	1	Yes	Private	Urban	
1	Female	61.0	0	0	Yes	Self-employed	Rural	
2	Male	80.0	0	1	Yes	Private	Rural	
3	Female	49.0	0	0	Yes	Private	Urban	
4	Female	79.0	1	0	Yes	Self-employed	Rural	
...	...	...	...	...	...	...	...	...
5105	Female	80.0	1	0	Yes	Private	Urban	
5106	Female	81.0	0	0	Yes	Self-employed	Urban	
5107	Female	35.0	0	0	Yes	Self-employed	Rural	
5108	Male	51.0	0	0	Yes	Private	Rural	
5109	Female	44.0	0	0	Yes	Govt_job	Urban	

5110 rows × 11 columns



In [6]:

```
df.shape
```

Out[6]: (5110, 11)

In [7]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   gender          5110 non-null   object  
 1   age              5110 non-null   float64 
 2   hypertension     5110 non-null   int64   
 3   heart_disease   5110 non-null   int64   
 4   ever_married    5110 non-null   object  
 5   work_type        5110 non-null   object  
 6   Residence_type  5110 non-null   object  
 7   avg_glucose_level 5110 non-null   float64 
 8   bmi              4909 non-null   float64 
 9   smoking_status  5110 non-null   object  
 10  stroke           5110 non-null   int64  
dtypes: float64(3), int64(3), object(5)
memory usage: 439.3+ KB
```

In [8]: `df.describe()`

	age	hypertension	heart_disease	avg_glucose_level	bmi	stroke
count	5110.000000	5110.000000	5110.000000	5110.000000	4909.000000	5110.000000
mean	43.226614	0.097456	0.054012	106.147677	28.893237	0.048728
std	22.612647	0.296607	0.226063	45.283560	7.854067	0.215320
min	0.080000	0.000000	0.000000	55.120000	10.300000	0.000000
25%	25.000000	0.000000	0.000000	77.245000	23.500000	0.000000
50%	45.000000	0.000000	0.000000	91.885000	28.100000	0.000000
75%	61.000000	0.000000	0.000000	114.090000	33.100000	0.000000
max	82.000000	1.000000	1.000000	271.740000	97.600000	1.000000

In [9]: `df.isnull().sum()`

```
Out[9]: gender          0
        age            0
        hypertension    0
        heart_disease   0
        ever_married    0
        work_type       0
        Residence_type  0
        avg_glucose_level 0
        bmi             201
        smoking_status   0
        stroke          0
        dtype: int64
```

In [10]: `df.isnull().sum().sort_values(ascending=False)`

```
Out[10]: bmi           201
        gender         0
        age            0
        hypertension    0
        heart_disease   0
        ever_married    0
        work_type       0
        Residence_type  0
        avg_glucose_level 0
        smoking_status   0
        stroke          0
        dtype: int64
```

```
In [11]: df.nunique()
```

```
Out[11]: gender           3  
age              104  
hypertension      2  
heart_disease     2  
ever_married       2  
work_type          5  
Residence_type     2  
avg_glucose_level 3979  
bmi               418  
smoking_status      4  
stroke              2  
dtype: int64
```

```
In [12]: df.nunique().sort_values(ascending=False)
```

```
Out[12]: avg_glucose_level    3979  
bmi                  418  
age                  104  
work_type             5  
smoking_status        4  
gender                 3  
hypertension            2  
heart_disease           2  
ever_married            2  
Residence_type           2  
stroke                  2  
dtype: int64
```

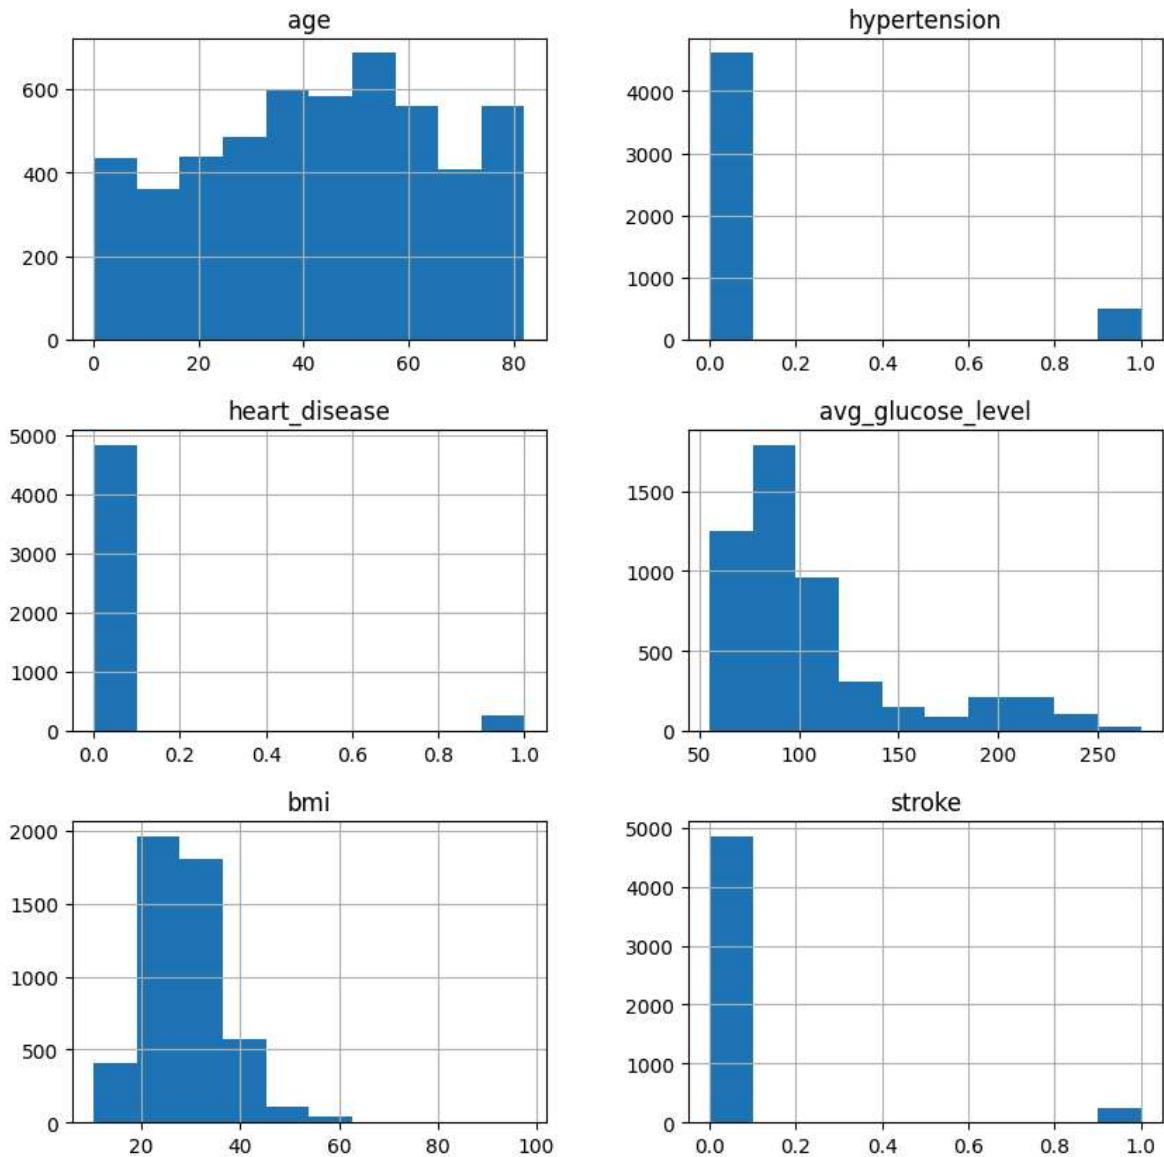
```
In [13]: df.duplicated().sum()
```

```
Out[13]: 0
```

```
In [14]: import matplotlib.pyplot as plt
```

```
In [15]: df.hist(figsize=(10,10))
plt.suptitle('Health Data Distribution: Histogram Analysis', fontsize=16)
plt.show()
```

### Health Data Distribution: Histogram Analysis



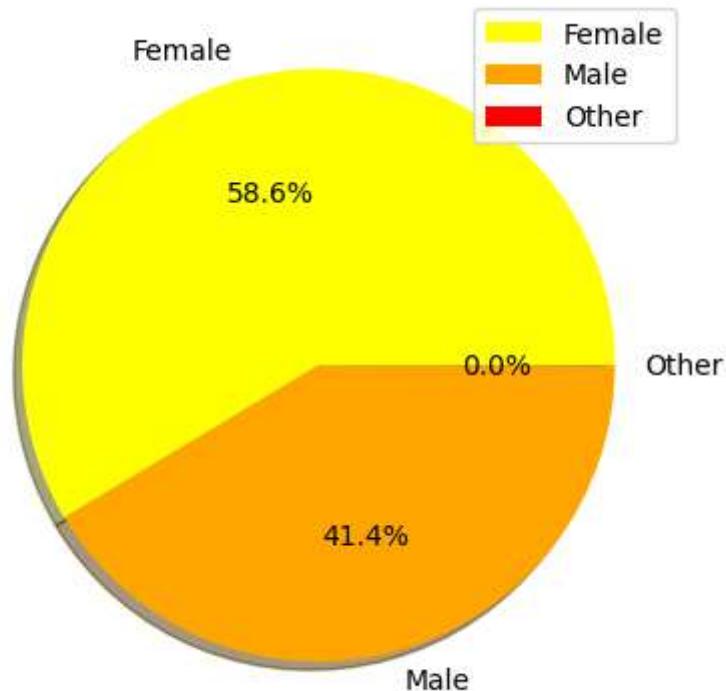
## Gender Analysis

```
In [17]: gender_count=df['gender'].value_counts()  
gender_percent=df['gender'].value_counts(normalize=True)*100  
table1=pd.DataFrame({'Count':gender_count,'Percentage %':gender_percent})  
table1
```

Out[17]:

gender	Count	Percentage %
Female	2994	58.590998
Male	2115	41.389432
Other	1	0.019569

```
In [18]: gender_counts = df['gender'].value_counts()  
plt.pie(gender_counts, labels=gender_counts.index, colors=['yellow', 'orange',  
plt.legend()  
plt.show()
```

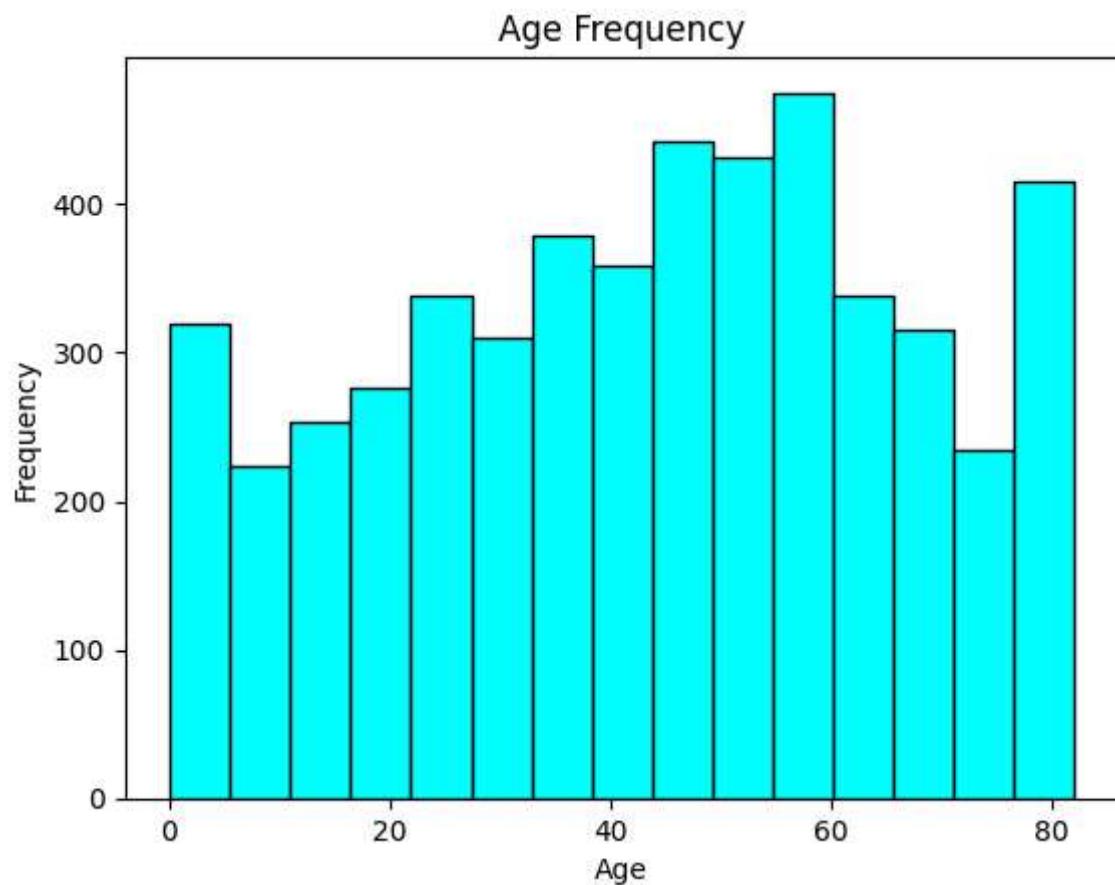


## Age Analysis

```
In [20]: df['age'].describe()
```

```
Out[20]: count    5110.000000
mean      43.226614
std       22.612647
min       0.080000
25%      25.000000
50%      45.000000
75%      61.000000
max      82.000000
Name: age, dtype: float64
```

```
In [21]: plt.hist(df['age'],bins=15,color='cyan',edgecolor='black',linewidth=1)
plt.title('Age Frequency')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



In [22]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5110 entries, 0 to 5109
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   gender          5110 non-null    object  
 1   age              5110 non-null    float64 
 2   hypertension     5110 non-null    int64  
 3   heart_disease   5110 non-null    int64  
 4   ever_married    5110 non-null    object  
 5   work_type        5110 non-null    object  
 6   Residence_type  5110 non-null    object  
 7   avg_glucose_level 5110 non-null    float64 
 8   bmi              4909 non-null    float64 
 9   smoking_status  5110 non-null    object  
 10  stroke           5110 non-null    int64  
dtypes: float64(3), int64(3), object(5)
memory usage: 439.3+ KB
```

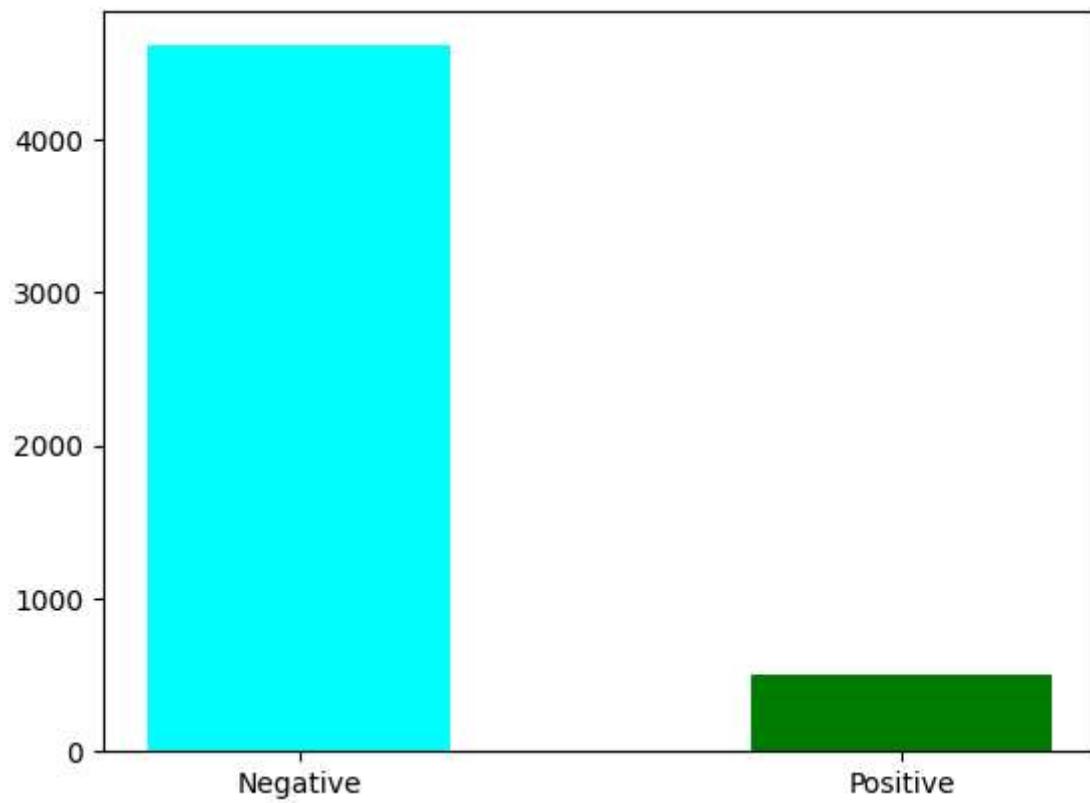
## Hyper Tension Analysis

In [24]: hyperTension\_count=df['hypertension'].value\_counts()  
hyperTension\_percent=df['hypertension'].value\_counts(normalize=True)\*100  
table2=pd.DataFrame({'Count':hyperTension\_count,'Percentage %':hyperTension\_percent})  
table2

Out[24]:

	Count	Percentage %
<b>hypertension</b>		
0	4612	90.254403
1	498	9.745597

```
In [25]: plt.bar(['Negative','Positive'],df['hypertension'].value_counts(),align='center')
plt.show()
```



## Heart Disease Analysis

```
In [27]: heartDisease_count=df['heart_disease'].value_counts()
heartDisease_percent=df['heart_disease'].value_counts(normalize=True)*100
table3=pd.DataFrame({'Count':heartDisease_count,'Percentage %':heartDisease_percent})
table3
```

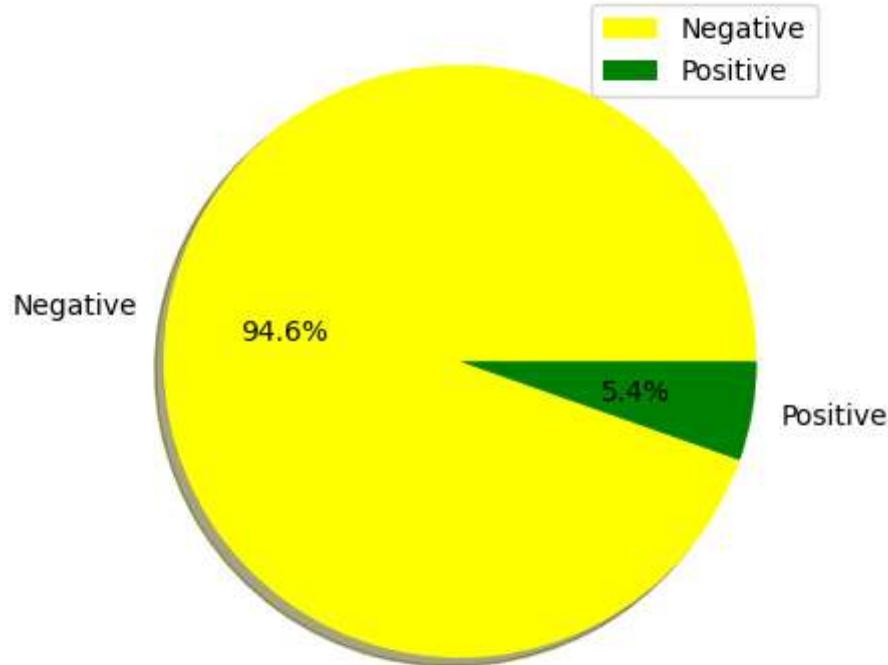
Out[27]:

heart_disease	Count	Percentage %
0	4834	94.598826
1	276	5.401174

heart\_disease

0	4834	94.598826
1	276	5.401174

```
In [28]: plt.pie(df['heart_disease'].value_counts(), labels=['Negative', 'Positive'], colors=[#FFFF00, #008000], autopct='%1.1f%%', startangle=90)
```



## Marriage Analysis

```
In [30]: df[['ever_married']]
```

```
Out[30]: ever_married
```

	ever_married
0	Yes
1	Yes
2	Yes
3	Yes
4	Yes
...	...
5105	Yes
5106	Yes
5107	Yes
5108	Yes
5109	Yes

5110 rows × 1 columns

```
In [31]: married_count=df['ever_married'].value_counts()  
married_percent=df['ever_married'].value_counts(normalize=True)*100  
table4=pd.DataFrame({'Count':married_count, 'Percentage %':married_percent})  
table4
```

Out[31]:

	Count	Percentage %
<b>ever_married</b>		
<b>Yes</b>	3353	65.616438
<b>No</b>	1757	34.383562

	Count	Percentage %
<b>ever_married</b>		
<b>Yes</b>	3353	65.616438
<b>No</b>	1757	34.383562

```
In [32]: import seaborn as sns
```

```
In [33]: sns.countplot(x='ever_married', palette=sns.color_palette("Paired"), data=df)
```

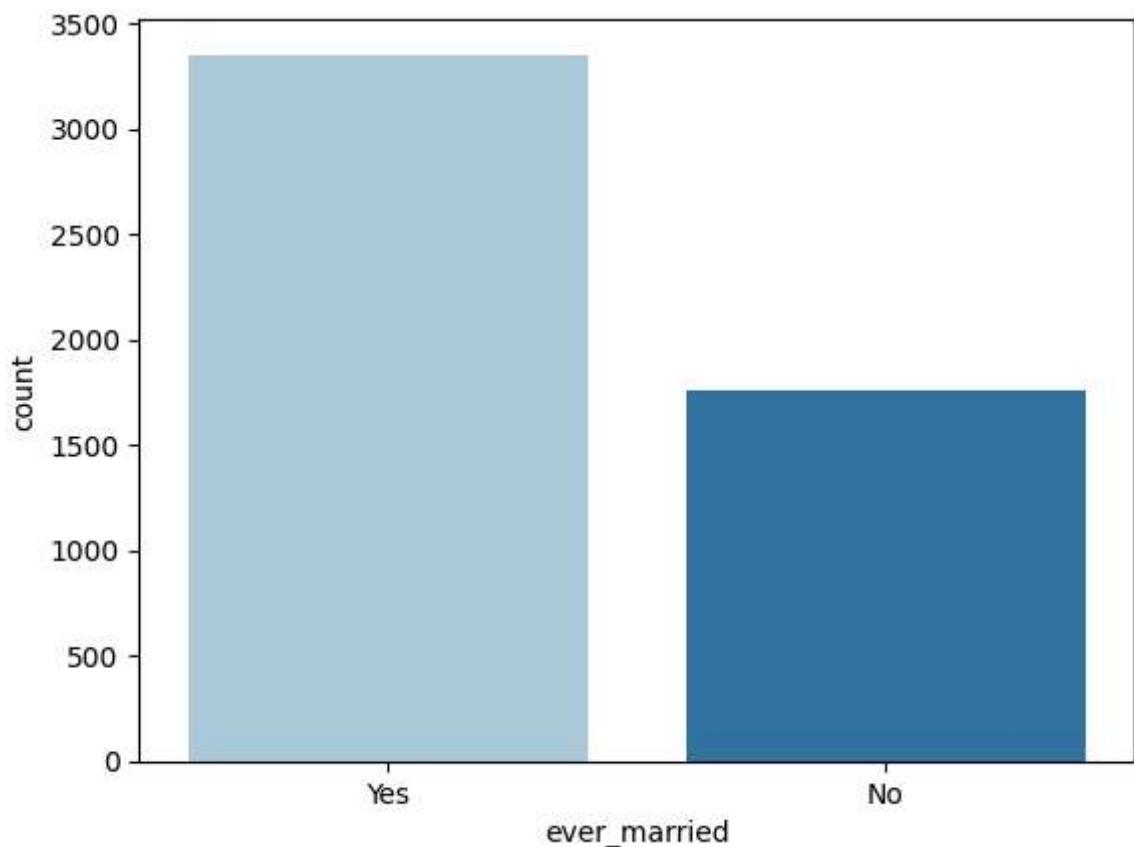
C:\Users\Udayabhanu\AppData\Local\Temp\ipykernel\_12612\3970370607.py:1: FutureWarning:  
Passing `palette` without assigning `hue` is deprecated and will be removed  
in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the  
same effect.

```
sns.countplot(x='ever_married', palette=sns.color_palette("Paired"), data=df)
```

C:\Users\Udayabhanu\AppData\Local\Temp\ipykernel\_12612\3970370607.py:1: UserWarning:  
The palette list has more values (12) than needed (2), which may not be intended.

```
sns.countplot(x='ever_married', palette=sns.color_palette("Paired"), data=df)
```

Out[33]: <Axes: xlabel='ever\_married', ylabel='count'>



## Work Type Analysis

```
In [35]: job_count=df['work_type'].value_counts()  
job_percent=df['work_type'].value_counts(normalize=True)*100  
table5=pd.DataFrame({'Count':job_count,'Percentage %':job_percent})  
table5
```

Out[35]:

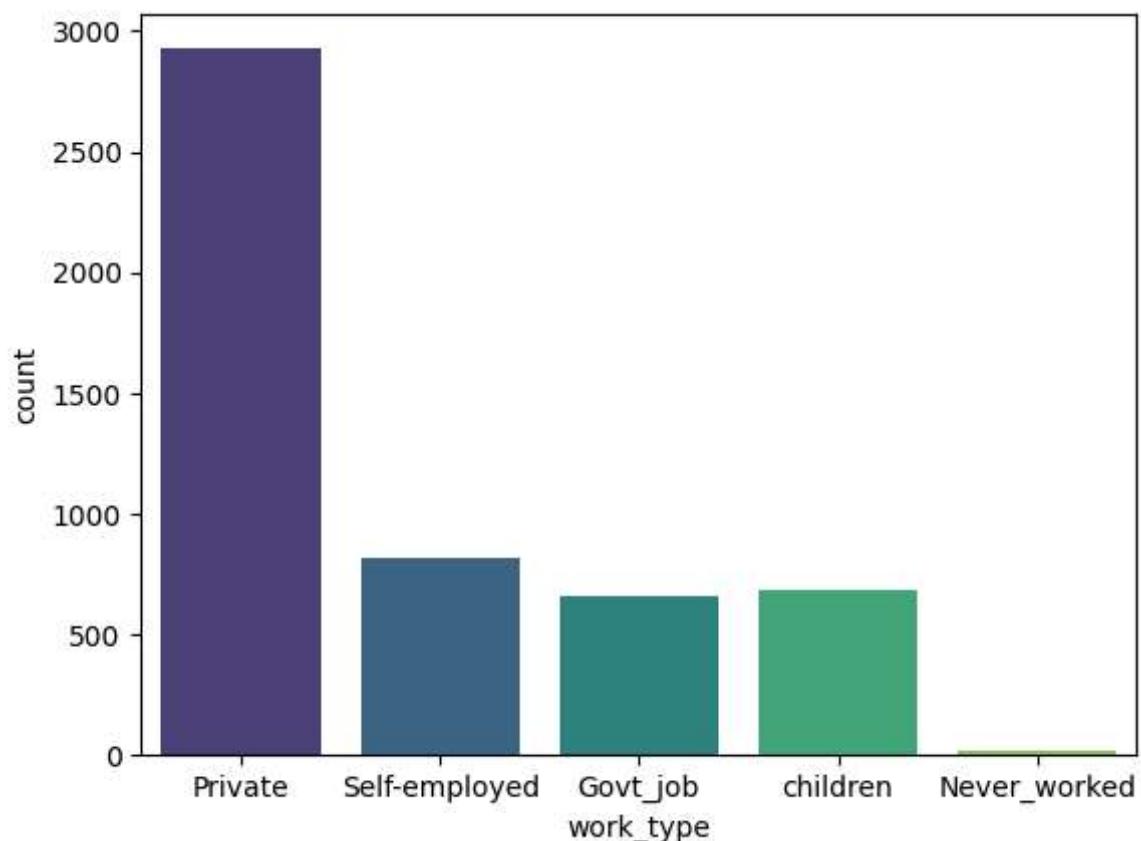
work_type	Count	Percentage %
Private	2925	57.240705
Self-employed	819	16.027397
children	687	13.444227
Govt_job	657	12.857143
Never_worked	22	0.430528

In [36]: `sns.countplot(x='work_type', palette='viridis', data=df)`

```
C:\Users\Udayabhanu\AppData\Local\Temp\ipykernel_12612\3879642141.py:1: FutureWarning:
Passing `palette` without assigning `hue` is deprecated and will be removed
in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the
same effect.
```

```
sns.countplot(x='work_type', palette='viridis', data=df)
```

Out[36]: <Axes: xlabel='work\_type', ylabel='count'>



## Residence Type Analysis

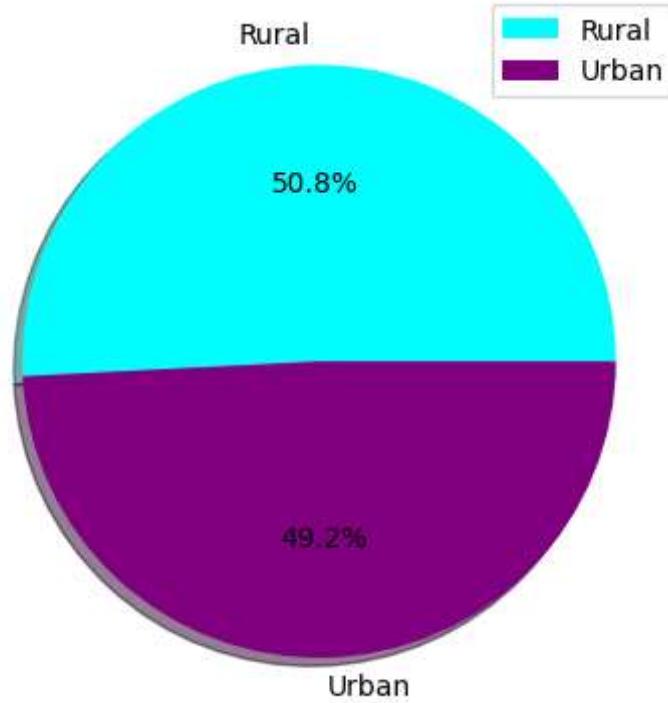
In [38]: `Residence_type_count=df['Residence_type'].value_counts()  
Residence_type_percent=df['Residence_type'].value_counts(normalize=True)*100  
table6=pd.DataFrame({'Count':Residence_type_count,'Percentage %':Residence_type_percent})  
table6`

Out[38]:

Residence_type	Count	Percentage %
Urban	2596	50.802348
Rural	2514	49.197652

Residence_type	Count	Percentage %
Urban	2596	50.802348
Rural	2514	49.197652

```
In [39]: plt.pie(df['Residence_type'].value_counts(), labels=['Rural','Urban'], colors=[plt.legend()  
plt.show())
```

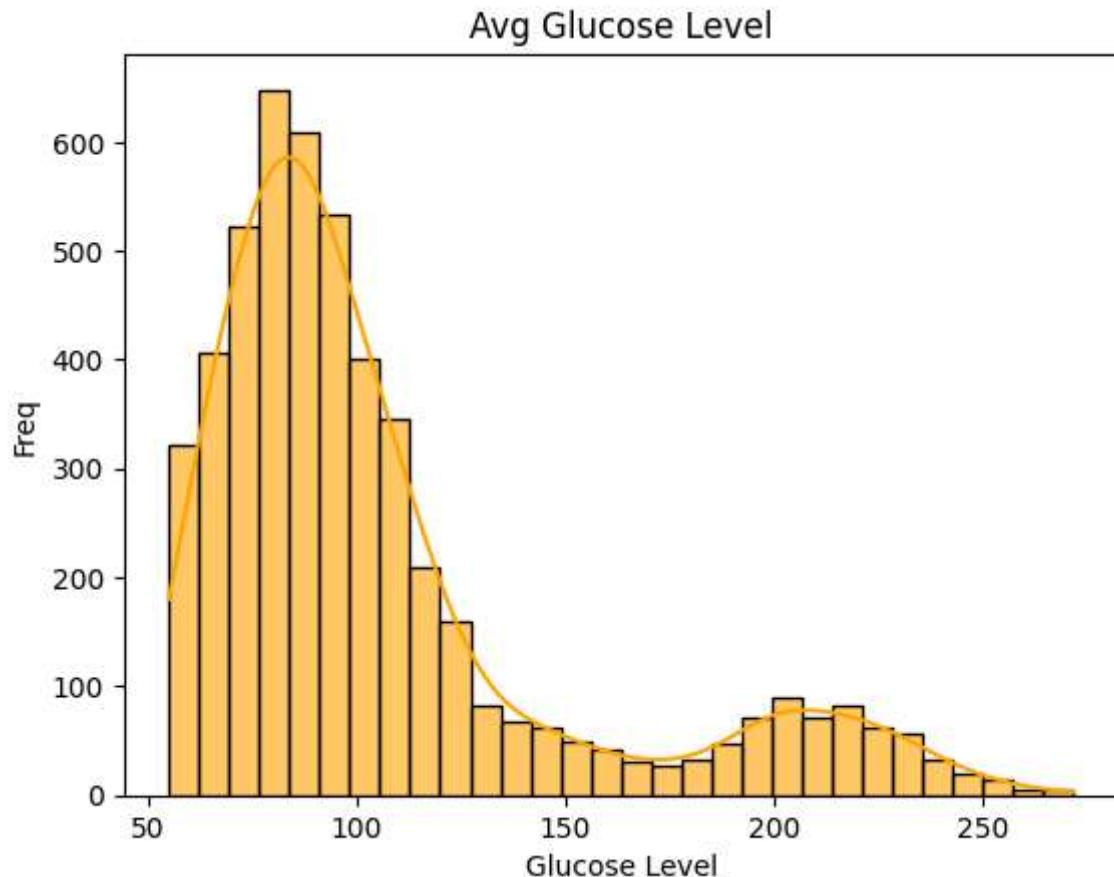


## Average Glucose Level Analysis

```
In [41]: df['avg_glucose_level'].describe()
```

```
Out[41]: count    5110.000000  
mean      106.147677  
std       45.283560  
min       55.120000  
25%       77.245000  
50%       91.885000  
75%      114.090000  
max      271.740000  
Name: avg_glucose_level, dtype: float64
```

```
In [42]: sns.histplot(df['avg_glucose_level'], bins=30, kde=True, color='orange', edgecolor='black')
plt.title('Avg Glucose Level')
plt.xlabel('Glucose Level')
plt.ylabel('Freq')
plt.show()
```



## Body Mass Index Analysis

```
In [44]: df['bmi'].describe()
```

```
Out[44]: count    4909.000000
          mean     28.893237
          std      7.854067
          min     10.300000
          25%    23.500000
          50%    28.100000
          75%    33.100000
          max    97.600000
          Name: bmi, dtype: float64
```

```
In [45]: df['bmi'].isnull().sum()
```

```
Out[45]: 201
```

```
In [46]: df.dropna(inplace=True)
```

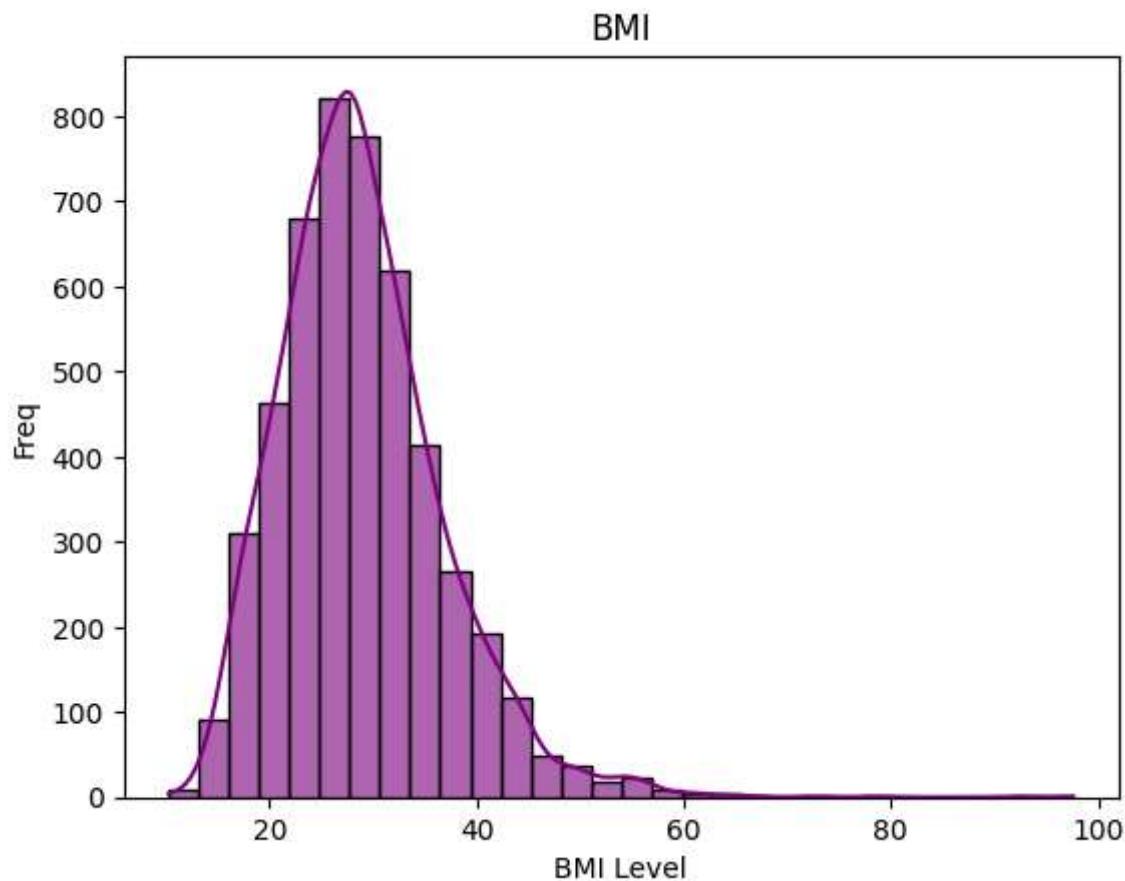
```
In [47]: df[df['bmi'].isna()]
```

```
Out[47]: gender age hypertension heart_disease ever_married work_type Residence_type avg_gluc
```

```
In [48]: df['bmi'].describe()
```

```
Out[48]: count    4909.000000
mean      28.893237
std       7.854067
min      10.300000
25%     23.500000
50%     28.100000
75%     33.100000
max      97.600000
Name: bmi, dtype: float64
```

```
In [49]: sns.histplot(df['bmi'], bins=30, kde=True, color='purple', edgecolor='black', line_kde=True)
plt.title('BMI')
plt.xlabel('BMI Level')
plt.ylabel('Freq')
plt.show()
```



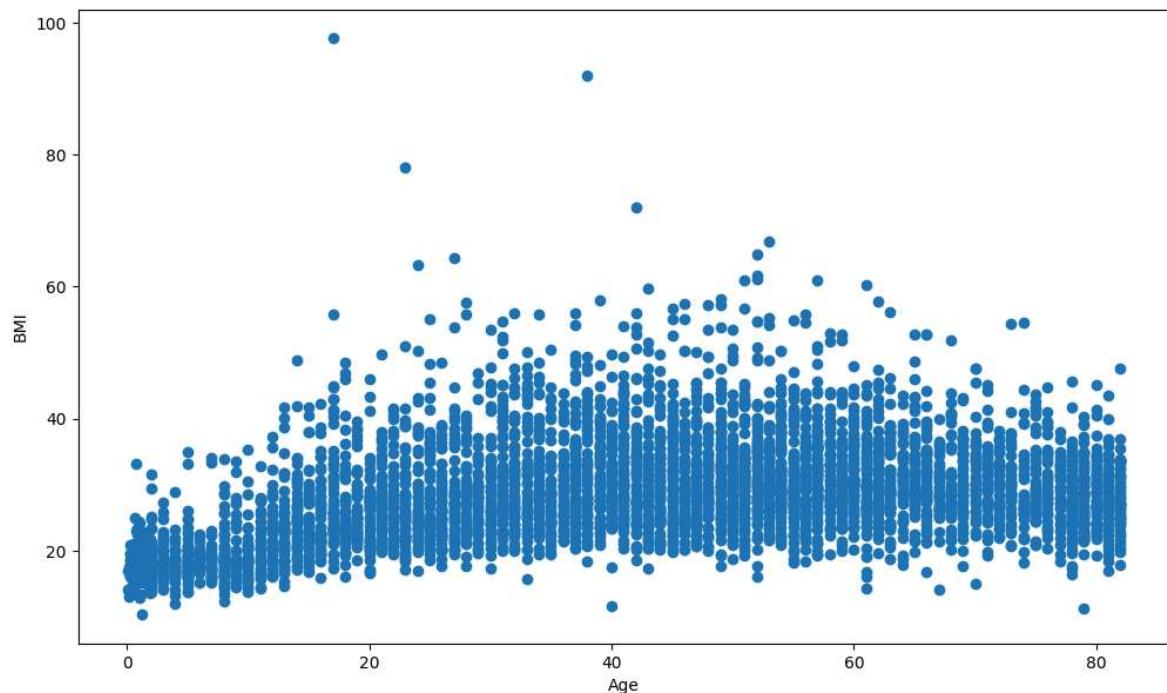
## Stroke Prediction Analysis

```
In [51]: stroke_count=df['stroke'].value_counts()
stroke_percent=df['stroke'].value_counts(normalize=True)*100
table7=pd.DataFrame({'Count':stroke_count,'Percentage %':stroke_percent})
table7
```

Out[51]:

stroke	Count	Percentage %
0	4700	95.742514
1	209	4.257486

```
In [52]: plt.figure(figsize=(12,7))
plt.scatter(x='age',y='bmi',data=df)
plt.xlabel('Age')
plt.ylabel('BMI')
plt.show()
```

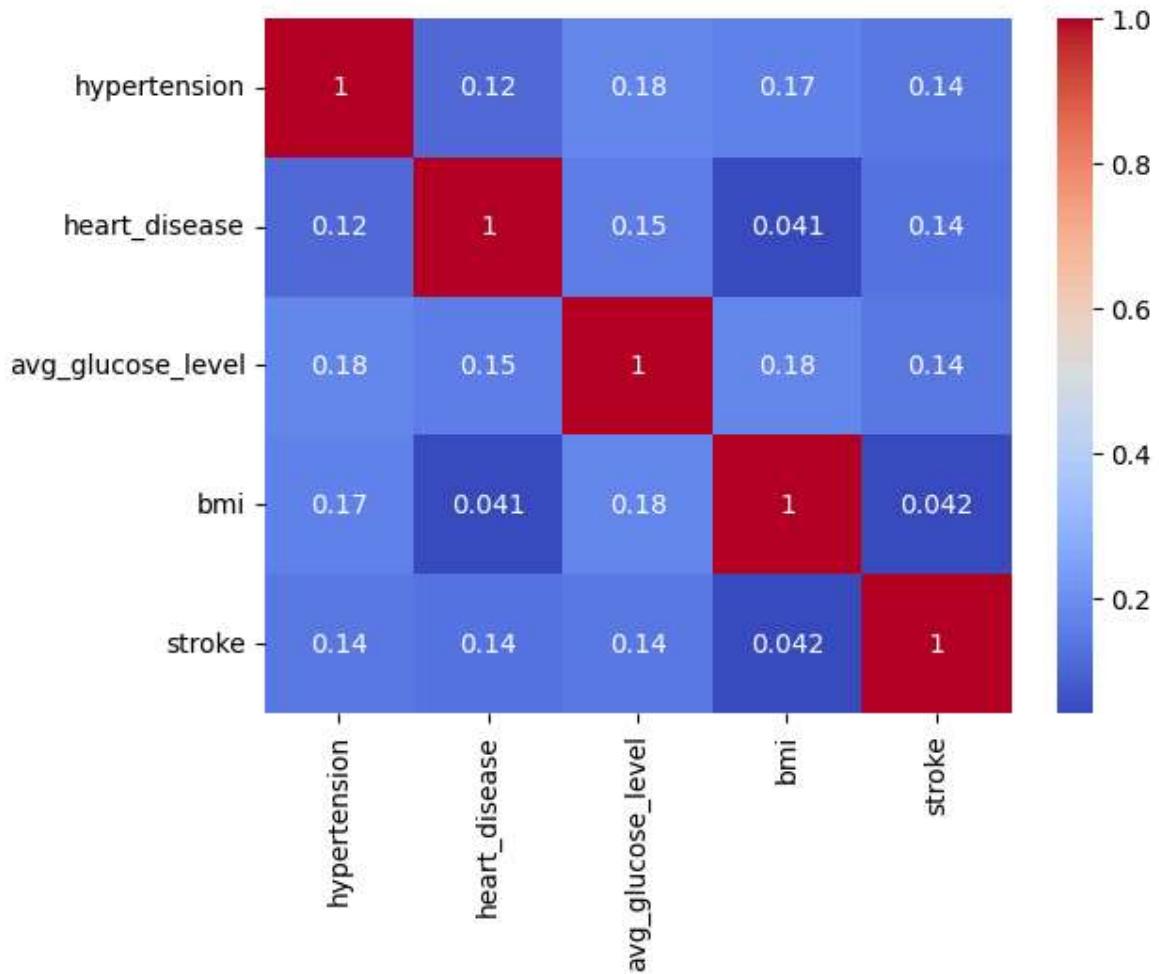


```
In [53]: df.groupby(['Residence_type'])['stroke'].value_counts()
```

```
Out[53]: Residence_type  stroke
Rural                 0      2319
                      1       100
Urban                0      2381
                      1       109
Name: count, dtype: int64
```

```
In [54]: df2=df.iloc[:, 2:]
cm = df2.corr(numeric_only=True)
sns.heatmap(cm, annot=True, cmap='coolwarm')
```

Out[54]: <Axes: >



```
In [ ]:
```

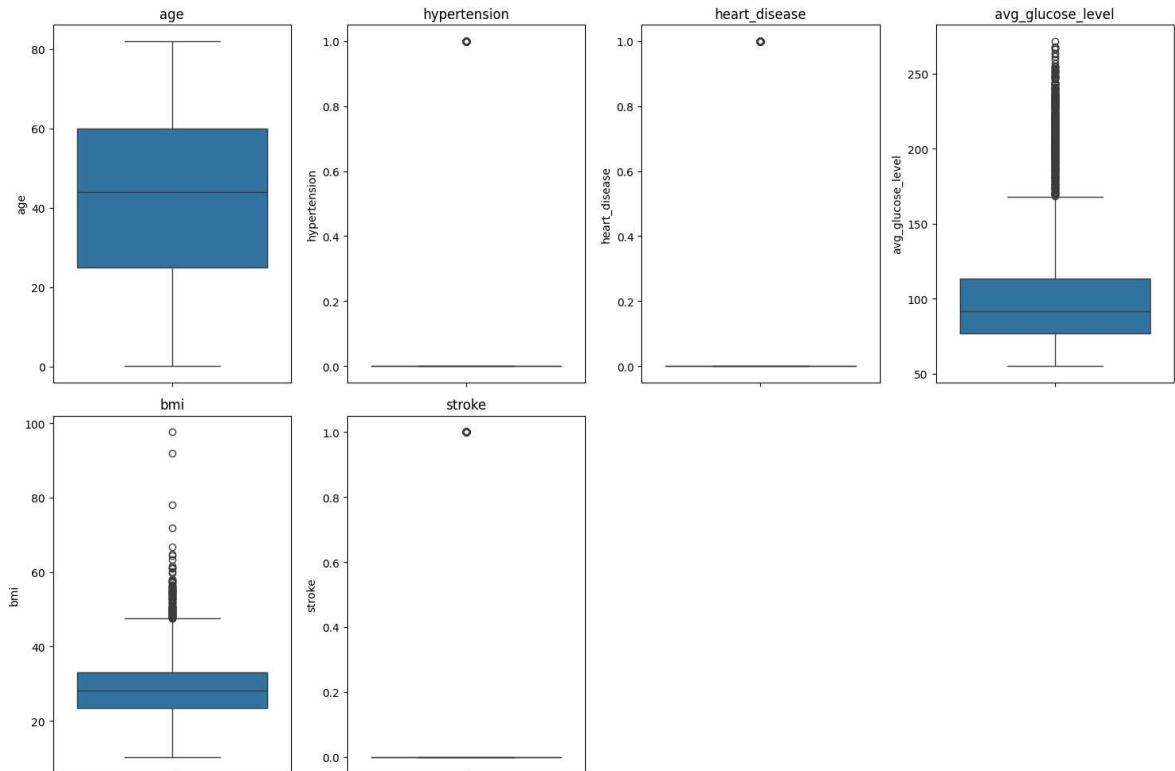
```
In [55]: numerical_columns = df.select_dtypes(include=['int64', 'float64']).columns
n_cols = 4
n_rows = len(numerical_columns) // n_cols + 1

plt.figure(figsize=(15, 5 * n_rows))

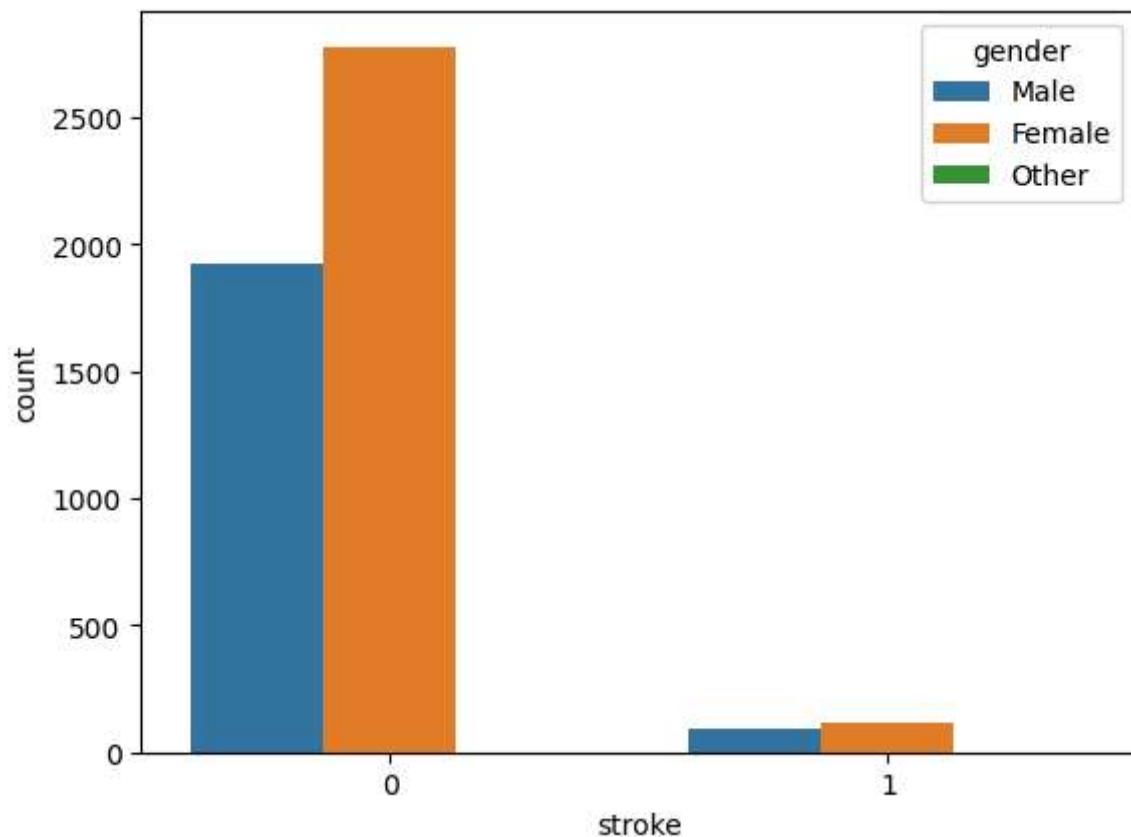
# Plot each numerical column as a boxplot
for i, col_name in enumerate(numerical_columns, 1):
    plt.subplot(n_rows, n_cols, i)
    sns.boxplot(y=df[col_name])
    plt.title(col_name)

# Adjust the layout for better spacing
plt.tight_layout()

# Display the plot
plt.show()
```

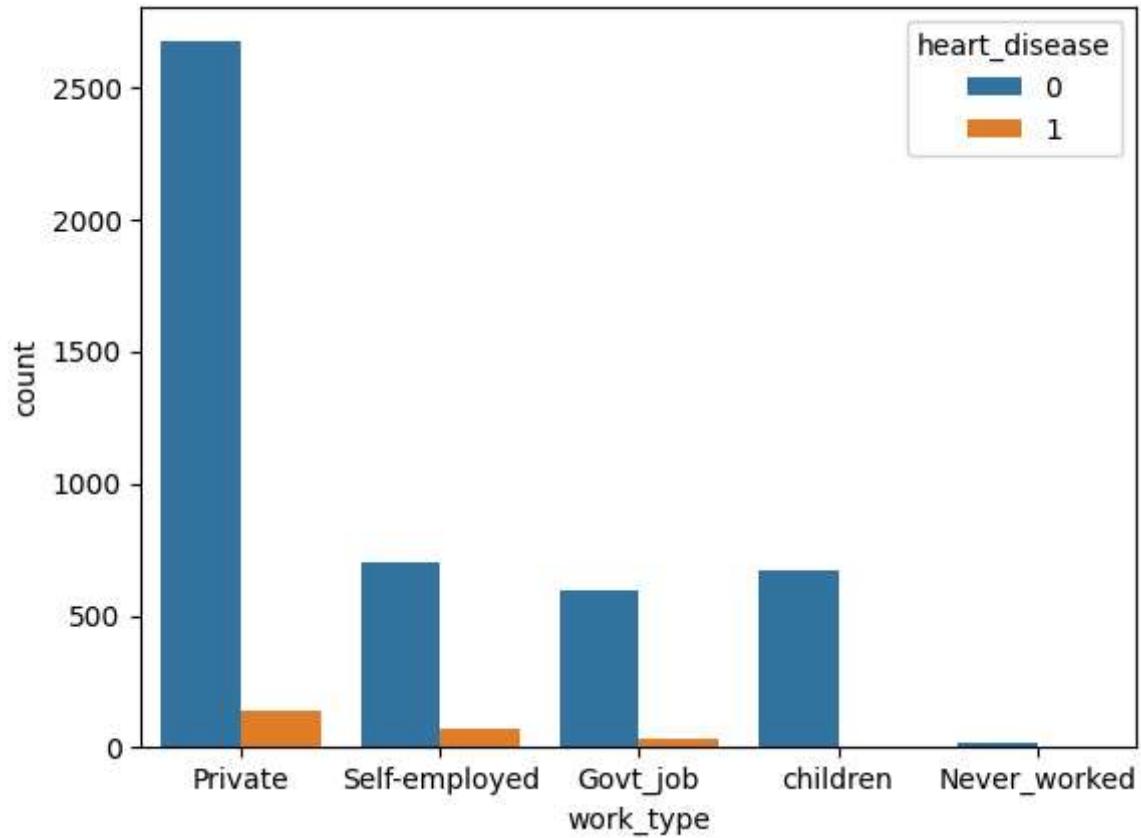


```
In [56]: sns.countplot(x='stroke',hue='gender', data=df)  
plt.show()
```



```
In [57]: sns.countplot(x = 'work_type',hue='heart_disease' ,data=df)
```

```
Out[57]: <Axes: xlabel='work_type', ylabel='count'>
```



## Correlations

```
In [59]: # Select columns of type int or float
obj_cols = df.select_dtypes(include=['int64', 'float64']).columns

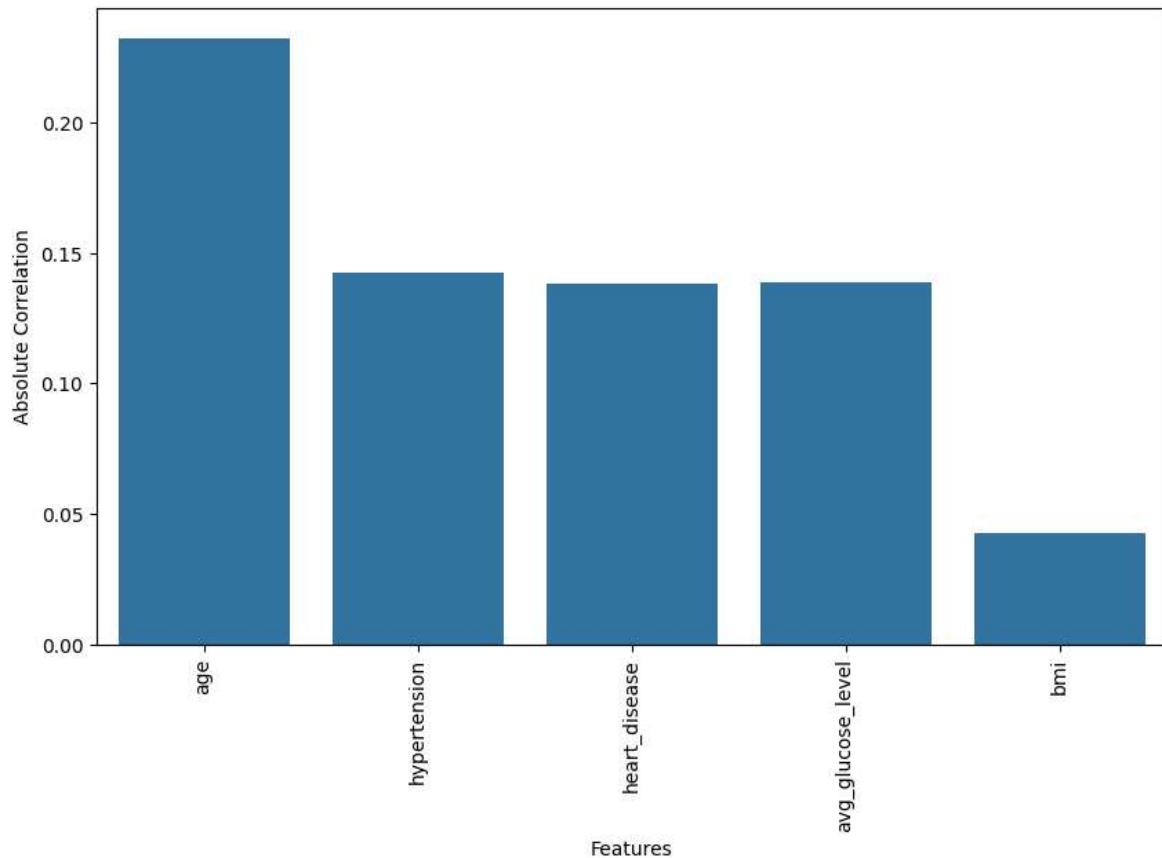
key = []
vals = []

# Iterate over numeric columns and compute correlation with 'RiskScore'
for n in obj_cols:
    if n=='stroke':
        continue
    key.append(n)
    vals.append(abs(df['stroke'].corr(df[n])))
    print(n,": ",df['stroke'].corr(df[n]))

# Plot the values
plt.figure(figsize=(10,6))
sns.barplot(x=key, y=vals)
plt.xticks(rotation=90) # Rotate x-axis Labels if necessary
plt.title('Correlation of Features with stroke')
plt.ylabel('Absolute Correlation')
plt.xlabel('Features')
plt.show()
```

```
age : 0.23233085553484958
hypertension : 0.1425146058811262
heart_disease : 0.13793778762219294
avg_glucose_level : 0.13893586200507035
bmi : 0.04237366114923355
```

## Correlation of Features with stroke



```
In [60]: health_data = pd.read_csv(r'C:\Users\Udayabhanu\Downloads\Stroke Prediction\Stro
```

```
In [61]: health_data = health_data.drop('id',axis=1)
```

```
In [62]: from sklearn.impute import SimpleImputer
impute= SimpleImputer(missing_values=np.nan,strategy='median')

health_data['bmi']=pd.DataFrame(impute.fit_transform(np.array(health_data['bmi'].values)))
health_data.isnull().sum()
```

```
Out[62]: gender          0
age             0
hypertension     0
heart_disease    0
ever_married      0
work_type         0
Residence_type    0
avg_glucose_level 0
bmi              0
smoking_status     0
stroke            0
dtype: int64
```

```
In [63]: from sklearn.preprocessing import OneHotEncoder
OH=OneHotEncoder()

for col in health_data.select_dtypes(include=['object']):
    health_data[col]=OH.fit_transform(health_data[[col]]).toarray()
health_data.head()
```

Out[63]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glu
0	0.0	67.0	0	1	0.0	0.0	0.0	0.0
1	1.0	61.0	0	0	0.0	0.0	0.0	1.0
2	0.0	80.0	0	1	0.0	0.0	0.0	1.0
3	1.0	49.0	0	0	0.0	0.0	0.0	0.0
4	1.0	79.0	1	0	0.0	0.0	0.0	1.0

```
In [64]: x=health_data.drop('stroke',axis=1)
y=health_data['stroke']
```

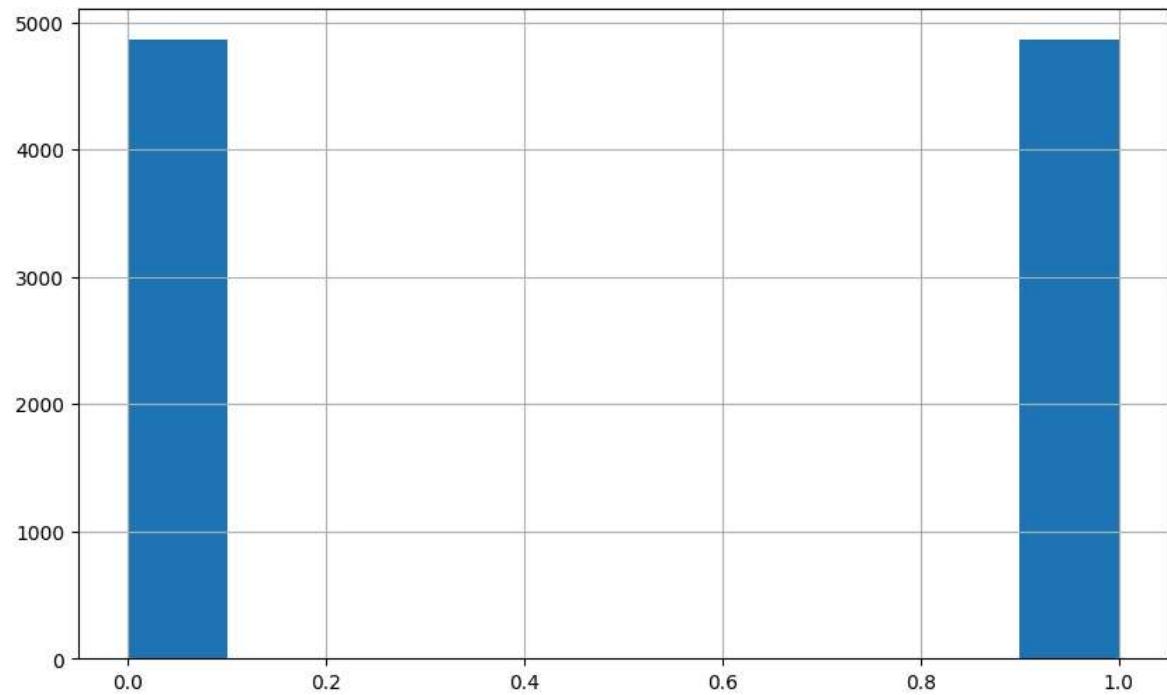
## Handling unbalanced data

```
In [66]: from imblearn.over_sampling import SMOTE
smote=SMOTE()

x_smote,y_smote=smote.fit_resample(x,y)
```

```
In [67]: plt.figure(figsize=(10, 6))
y_smote.hist()
```

```
Out[67]: <Axes: >
```



### Robust scaler

```
In [69]: from sklearn.preprocessing import RobustScaler
rs=RobustScaler()

x_smote=rs.fit_transform(x_smote)
```

```
In [70]: from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test=train_test_split(x_smote,y_smote,test_size=0.2,
```

# Machine Learning

```
In [72]: from sklearn.metrics import accuracy_score, f1_score, precision_score, recall
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC

# Dictionary of models
models = {

    'LogisticRegression': LogisticRegression(max_iter = 2000),

    'RandomForestClassifier': RandomForestClassifier(),

    'KNeighborsClassifier' : KNeighborsClassifier(),

    'DecisionTreeClassifier': DecisionTreeClassifier(),

    'GaussianNB' : GaussianNB(),

    'Support Vector Machine' : SVC() }
```

```
In [73]: from tqdm import tqdm
# Fit models, predict and calculate accuracy and F1 score
results = []
models_name = []
for name, model in tqdm(models.items()):
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    accuracy = accuracy_score(y_test, predictions)
    f1 = f1_score(y_test, predictions, average='weighted')
    precision = precision_score(y_test, predictions, average='weighted')
    recall = recall_score(y_test, predictions, average='weighted')
    models_name.append(name)
    results.append([accuracy,precision,recall,f1])
```

100% |  | 6/6 [00:02<00:00, 2.42it/s]

```
In [74]: Model_accuracy = pd.DataFrame(results, index=models_name, columns = ['Accuracy'])
```

Model\_accuracy

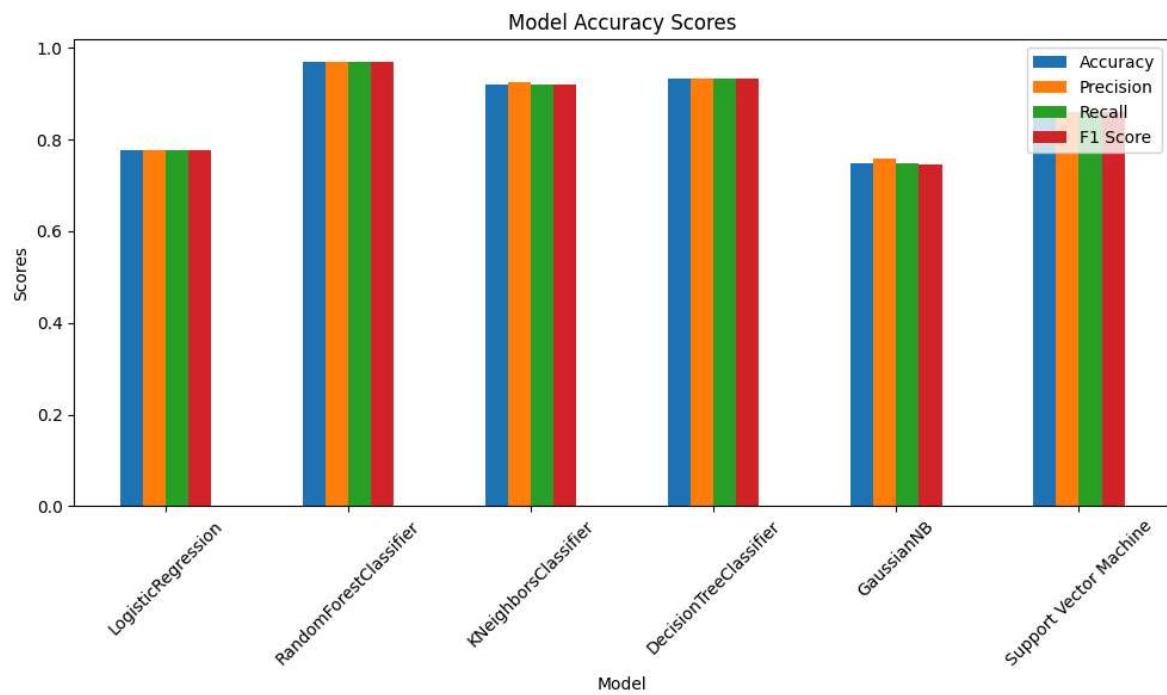
Out[74]:

	Accuracy	Precision	Recall	F1 Score
<b>LogisticRegression</b>	0.776350	0.777269	0.776350	0.776187
<b>RandomForestClassifier</b>	0.969152	0.969596	0.969152	0.969143
<b>KNeighborsClassifier</b>	0.919280	0.924245	0.919280	0.919056
<b>DecisionTreeClassifier</b>	0.933676	0.934123	0.933676	0.933662
<b>GaussianNB</b>	0.749100	0.758677	0.749100	0.746848
<b>Support Vector Machine</b>	0.860154	0.860258	0.860154	0.860148

```
In [75]: Model_accuracy.plot(kind='bar', figsize=(10, 6))
```

```
# Customizing the plot
plt.xlabel('Model')
plt.ylabel('Scores')
plt.title('Model Accuracy Scores')
plt.xticks(rotation=45) # Rotate model names for better readability
plt.legend(loc='upper right')
plt.tight_layout() # Adjust layout to fit labels

# Display the plot
plt.show()
```



## Neural Network

```
In [77]: import tensorflow as tf  
from tensorflow import keras
```

```
In [78]: model = keras.Sequential()  
model.add(keras.layers.Dense(50, activation="relu", input_shape = X_train.shape[1]))  
model.add(keras.layers.Dense(40, activation="relu"))  
model.add(keras.layers.Dense(30, activation="relu"))  
model.add(keras.layers.Dense(20, activation="relu"))  
model.add(keras.layers.Dense(10, activation="relu"))  
model.add(keras.layers.Dense(1, activation="sigmoid"))
```

```
C:\Users\Udayabhanu\AppData\Roaming\Python\Python312\site-packages\keras\src  
\layers\core\dense.py:87: UserWarning: Do not pass an `input_shape`/`input_d  
im` argument to a layer. When using Sequential models, prefer using an `Inpu  
t(shape)` object as the first layer in the model instead.  
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [79]: model.compile(loss="binary_crossentropy", optimizer="adam", metrics=["accuracy"]  
  
# Training and evaluating the model  
history = model.fit(X_train, y_train, epochs=100, batch_size=100, validation_
```

```
Epoch 1/100  
63/63 2s 4ms/step - accuracy: 0.6875 - loss: 0.6108  
- val_accuracy: 0.7982 - val_loss: 0.4319  
Epoch 2/100  
63/63 0s 2ms/step - accuracy: 0.8061 - loss: 0.4293  
- val_accuracy: 0.8201 - val_loss: 0.3938  
Epoch 3/100  
63/63 0s 2ms/step - accuracy: 0.8261 - loss: 0.3899  
- val_accuracy: 0.8361 - val_loss: 0.3719  
Epoch 4/100  
63/63 0s 1ms/step - accuracy: 0.8471 - loss: 0.3492  
- val_accuracy: 0.8406 - val_loss: 0.3582  
Epoch 5/100  
63/63 0s 2ms/step - accuracy: 0.8692 - loss: 0.3098  
- val_accuracy: 0.8663 - val_loss: 0.3123  
Epoch 6/100  
63/63 0s 2ms/step - accuracy: 0.8741 - loss: 0.2887  
- val_accuracy: 0.8798 - val_loss: 0.2901  
Epoch 7/100  
63/63 0s 2ms/step - accuracy: 0.8807 - loss: 0.2562
```

In [80]: `model.predict(X_test)`

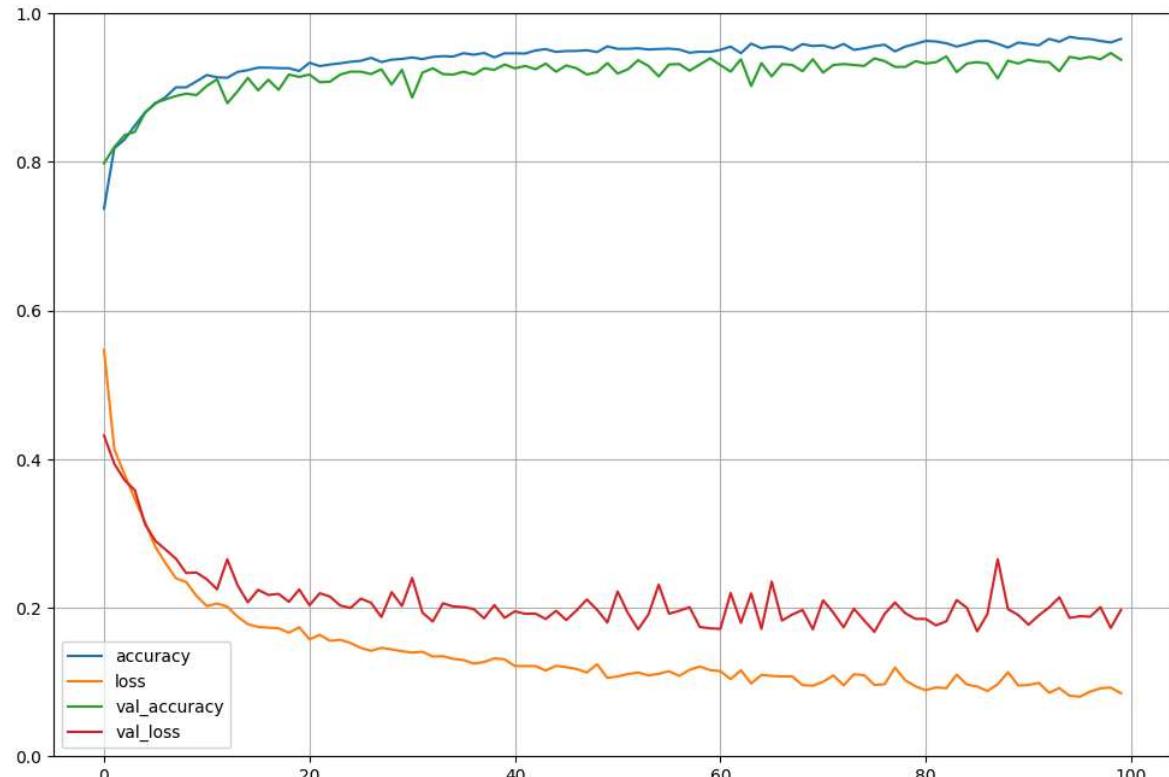
61/61 ━━━━━━━━ 0s 1ms/step

Out[80]: `array([[1.4419915e-07],  
[9.9999750e-01],  
[9.6922344e-01],  
...,  
[9.9995506e-01],  
[2.5440505e-08],  
[9.9983770e-01]], dtype=float32)`

In [81]: `import pandas as pd  
import matplotlib.pyplot as plt  
pd.DataFrame(history.history).plot(figsize=(12, 8))  
plt.grid(True)  
plt.gca().set_ylim(0, 1) # set the vertical range to [0-1]  
plt.show()`

*# Evaluate the model*

`model_evaluate = model.evaluate(X_test, y_test)  
print("Loss : ", model_evaluate[0])  
print("accuracy : ", model_evaluate[1])`



61/61 ━━━━━━━━ 0s 1ms/step - accuracy: 0.9223 - loss: 0.2063

Loss : 0.22870728373527527

accuracy : 0.9244216084480286

## Conclusion

**The project involved analyzing health data to predict stroke risk. Key achievements include:**

- Preprocessing data by handling missing values and encoding features.
- Conducting Exploratory Data Analysis (EDA) to visualize key health factors.
- Training machine learning models like logistic regression and random forests.
- Evaluating model performance and identifying important health indicators for stroke prediction.